

Capacity and Security of Heterogeneous Distributed Storage Systems

Toni Ernvall*, Salim El Rouayheb, *Member, IEEE*[†],
Camilla Hollanti, *Member, IEEE*[‡], and H. Vincent Poor, *Fellow, IEEE*[†]

*TUCS and Department of Mathematics and Statistics, University of Turku, Email: tmernv@utu.fi

[†]Department of Electrical Engineering, Princeton University, Emails: {salim, poor}@princeton.edu

[‡]Department of Mathematics and System Analysis, Aalto University, Finland, Email: camilla.hollanti@aalto.fi

Abstract—The capacity of heterogeneous distributed storage systems under repair dynamics is studied. Examples of these systems include peer-to-peer storage clouds, wireless, and Internet caching systems. Nodes in a heterogeneous system can have different storage capacities and different repair bandwidths. Lower and upper bounds on the system capacity are given. These bounds depend on either the average resources per node, or on a detailed knowledge of the node characteristics. Moreover, the case in which nodes may be compromised by an eavesdropper is addressed and bounds on the secrecy capacity of the system are derived. One implication of these new results is that symmetric repair maximizes the capacity of a homogeneous system, which justifies the model widely used in the literature.

I. INTRODUCTION

Cloud storage has emerged in recent years as an inexpensive and scalable solution for storing large amounts of data and making it pervasively available to users. The growing success of cloud storage has been accompanied by new advances in the theory of such systems, namely the application of network coding techniques for distributed data storage and the theory of regenerating codes introduced by Dimakis *et al.* [1], followed by a large body of further work in the literature.

Cloud storage systems are typically built using a large number of inexpensive commodity disks that fail frequently, making failures “the norm rather than the exception” [2]. Therefore, it is a prime concern to achieve fault-tolerance in these systems and minimize the probability of losing the stored data. The recent theoretical results uncovered fundamental tradeoffs among system resources (storage capacity, repair bandwidth, etc.) that are necessary to achieve fault-tolerance. They also provided novel code constructions for data redundancy schemes that can achieve these tradeoffs in certain cases; see for example [3] and the references within.

The majority of the results in the literature of this field focuses on a homogeneous model when studying the information theoretic limits on the performance of distributed storage systems. In a homogeneous system all the storage nodes have the same parameters (storage capacity, repair bandwidth, etc.). This model encompasses many real-world storage systems such as clusters in a data center, and has been instrumental in forming the engineering intuition for understanding these systems. Recent development have included the emergence of *heterogeneous* systems that pool together nodes from different sources and with different characteristics to form one big reliable cloud storage system. Examples include peer-to-peer

(p2p), or hybrid (p2p-assisted) cloud storage systems [4], [5], Internet caching systems for video-on-demand applications [6], [7], and caching systems in heterogeneous wireless networks [8]. Motivated by these applications, we study the capacity of heterogeneous distributed storage systems (DSSs) here under reliability and secrecy constraints.

Contributions: The capacity of a DSS is defined as the maximum amount of information that can be delivered to any user contacting k out of n nodes in the system. Intuitively, in a heterogeneous system, this capacity should be limited by the “weakest” nodes. However, nodes can have different storage capacities and different repair bandwidths. And the tension between these two set of parameters makes it challenging to identify which nodes are the “weakest”.

Our first result establishes an upper bound on the capacity of a DSS that depends on the average resources in the system (average storage capacity and average repair bandwidth per node). We use this bound to prove that symmetric repair, *i.e.*, downloading equal amount of data from each helper node, maximizes the capacity of a homogeneous DSS. While the optimality of symmetric repair is known for the special case of maximum distance separable (MDS) codes [9], our results assert that symmetric repair is always optimal and avoid the combinatorial cut-based arguments typically used in this context.

In addition, we give an expression for the capacity when we know the characteristics of all the nodes in the system (not just the averages). This expression may be hard to compute, but we use it to derive additional bounds that are easy to evaluate. Our techniques generalize to the scenario in which the system is compromised by an adversary. We focus on the case of a passive adversary (eavesdropper), and give bounds on the secrecy capacity when the system is supposed to leak no information (perfect secrecy). Moreover, we show that symmetric repair maximizes the secrecy capacity.

Related work: Wu proved the optimality of symmetric repair in [9] for the special case of systems using MDS codes. Codes for a non-homogeneous DSS with one super-node that is more reliable and has more storage capacity were studied in [10]. References [11] and [12] studied the problem of storage allocations in DSSs under a total storage budget constraint. Shah *et al.* studied in [13] constructions of regenerating codes that allow flexibility in splitting the repair bandwidth budget among the helper nodes. Pawar *et al.* [14] studied the secure capacity of DSSs under eavesdropping and malicious attacks.

Organization: The paper is organized as follows. In Section II, we describe our model for heterogeneous DSSs and set up the notation. In Section III, we summarize our main results. In Section IV, we prove our bounds on the capacity of a heterogeneous DSS. In Section V, we study the secrecy capacity in the presence of an eavesdropper. We conclude in Section VI and discuss some open problems.

II. MODEL

A heterogeneous distributed storage system is formed of n storage nodes v_1, \dots, v_n with storage capacities $\alpha_1, \dots, \alpha_n$ respectively. Unless stated otherwise, we assume that the nodes are indexed in increasing order of capacity, *i.e.*, $\alpha_1 \leq \alpha_2 \leq \dots \leq \alpha_n$. In a homogeneous system all nodes have the same storage capacity α , *i.e.*, $\alpha_i = \alpha, \forall i$. As a reliability requirement, a user should be able to obtain a file by contacting any $k < n$ nodes in the DSS. The nodes forming the system are unreliable and can fail. The system is *repaired* from a failure by replacing the failed node with a new one. Upon joining the system, the new node downloads its data from d , $k \leq d \leq n-1$, helper nodes in the system. We focus in this paper on *functional repair* in which the data stored on the replacement node need not be an exact copy of the lost data, but merely “functionally equivalent” [3]. Some of our results can be generalized to the exact repair case. See the extended version of this paper [15, Appendix A] for more details.

An important system parameter is the *repair bandwidth* which refers to the total amount of data downloaded by the new node. In a homogeneous system, the repair bandwidth, denoted by γ , is the same for any new node joining the system. The typical model adopted in the literature assumes *symmetric repair* in which the total repair bandwidth γ is divided equally among the d helpers. Thus, the new node downloads $\beta = \gamma/d$ amount of information from each helper. In a heterogeneous system the repair bandwidth can vary depending on which node has failed and which nodes are helping in the repair process. We denote by β_{ijS} the amount of information that a new node replacing the failed node v_j is downloading from node v_i when the other helper nodes belong to the index set S ($i \in S, |S| = d$). An important case is when the repair bandwidth per helper depends *only* on the the helper node. In this case, we say that helper node v_i has repair bandwidth β_i , *i.e.*, $\beta_{ijS} = \beta_i, \forall j, S$. In the case of a homogeneous system with symmetric repair, we have $\beta_{ijS} = \beta = \gamma/d, \forall i, j, S$.

We focus on repair from single node failures. In this case, there are $\binom{n-1}{d}$ possibilities for the set of helpers S . Therefore, the average repair bandwidth γ_j of node v_j is

$$\gamma_j = \binom{n-1}{d}^{-1} \sum_{\substack{S: j \notin S \\ |S|=d}} \sum_{i \in S} \beta_{ijS}. \quad (1)$$

We denote by $\bar{\gamma} = \frac{1}{n} \sum_{j=1}^n \gamma_j$ and $\bar{\alpha} = \frac{1}{n} \sum_{j=1}^n \alpha_j$ the average total repair bandwidth and average node capacity in the DSS, respectively. We are interested in finding the capacity C of a heterogeneous system. The capacity C represents the maximum amount of information that can be downloaded by any user contacting k out of the n nodes in the system. Recall

from [1], that the capacity C^{ho} of a homogeneous system implementing symmetric repair is given by

$$C^{ho}(\alpha, \gamma) = \sum_{i=1}^k \min \left\{ \alpha, (d-i+1) \frac{\gamma}{d} \right\}. \quad (2)$$

We are also interested in characterizing the secrecy capacity of the system when some nodes are compromised by an eavesdropper. We follow the model in [14] and denote by ℓ , $\ell \leq k$, the number of compromised nodes. The eavesdropper is assumed to be passive. She can read the data downloaded during repair and stored on a compromised node. We are interested here in information theoretic secrecy which characterizes the fundamental ability of the system to provide data confidentiality independently of cryptographic methods. The secrecy capacity of the system, denoted by C_s , is defined as the maximum amount of information that can be delivered to a user without revealing any information to the eavesdropper (perfect secrecy). We denote by C_s^{ho} the secrecy capacity of a homogeneous system with symmetric repair. Finding C_s^{ho} is still an open problem in general. The following upper bound was shown to hold in [14]:

$$C_s^{ho}(\alpha, \gamma, \ell) \leq \sum_{i=\ell+1}^k \min \left\{ \alpha, (d-i+1) \frac{\gamma}{d} \right\}. \quad (3)$$

III. MAIN RESULTS

We start by summarizing our results. Theorem 1 gives a general upper bound on the storage capacity of a heterogeneous DSS as a function of the average resources per node.

Theorem 1: The capacity C of a heterogeneous distributed storage system, with node average capacity $\bar{\alpha}$ and average repair bandwidth $\bar{\gamma}$, is upper bounded by

$$C \leq \sum_{i=1}^k \min \left\{ \bar{\alpha}, (d-i+1) \frac{\bar{\gamma}}{d} \right\} = C^{ho}(\bar{\alpha}, \bar{\gamma}). \quad (4)$$

The right-hand side term in (4) is the capacity of a homogeneous system in (2) in which all nodes have storage $\alpha = \bar{\alpha}$ and total repair bandwidth $\gamma = \bar{\gamma}$. Th. 1 states that the capacity of a DSS cannot exceed that of a homogeneous system where the total system resources are split equally among all the nodes. Moreover, Th. 1 implies that *symmetric repair is optimal* in homogeneous systems in the sense that it maximizes the system capacity. This justifies the repair model adopted in the literature. This result is stated formally in Cor. 2. While the optimality of symmetric repair is known for the special case of MDS codes [9], Cor. 2 asserts that symmetric repair is always optimal for any choice of system parameters. This result follows directly from Th. 1 and avoids the combinatorial arguments that may be needed in a direct proof.

Corollary 2: Symmetric repair maximizes the system capacity of a homogeneous DSS.

When we know the parameters of the nodes in the system beyond the averages, we can obtain possibly tighter capacity bounds as described in Th. 3. To simplify the notation, we order the repair bandwidth per helper β_{ijS} into an increasing




Fig. 1. An example that illustrates the proof of the capacity upper bound in (4). (a) A heterogeneous DSS with $(n, k, d) = (3, 2, 2)$ with node capacities $\alpha_1 = 1, \alpha_2 = \alpha_3 = 2$ and the repair bandwidth $\beta_1 = 1, \beta_2 = \beta_3 = 2$. (b) A DSS constructed by combining together $n! = 6$ copies of the original heterogeneous system corresponding to all possible node permutations. The obtained DSS is homogeneous with uniform storage per node $\alpha = 10$ and repair bandwidth per helper $\beta = 10$. The capacity of this system is 20 as given by (2) [1]. Any code that stores a file of size C ($C = 3$ here) on the original DSS can be transformed into a scheme that stores a file of size $n!C = 6C$ in the “bigger” system. This gives the upper bound in (4) $C \leq 20/6 = 10/3$.

sequence $\beta'_1, \beta'_2, \dots, \beta'_m$, such that $\beta'_l \leq \beta'_{l+1}$ and where $m = nd \binom{n-1}{d}$. Also, recall that $\alpha_1 \leq \alpha_2 \leq \dots \leq \alpha_n$.

Theorem 3: The capacity C of heterogeneous DSS is bounded by $C_{\min} \leq C \leq C_{\max}$, with

$$C_{\min} = \min_{l=0, \dots, k} \left(\sum_{j=1}^l \alpha_j + \sum_{j=1}^h \beta'_j \right),$$

$$C_{\max} = \min_{l=0, \dots, k} \left(\sum_{j=1}^l \alpha_j + \sum_{j=1}^h \beta'_{m+1-j} \right),$$

and $h = \frac{(2d-k-l+1)(k-l)}{2}$.

When the system is compromised by an eavesdropper the system secrecy capacity can be upper bounded as follows.

Theorem 4: The secrecy capacity C_s of a DSS when ℓ nodes are compromised by an eavesdropper is upper bounded by

$$C_s \leq \sum_{i=\ell+1}^k \min \left\{ \bar{\alpha}, (d-i+1) \frac{\bar{\gamma}}{d} \right\}. \quad (5)$$

This theorem implies that symmetric repair also maximizes the secrecy capacity of a homogeneous DSS.

IV. CAPACITY OF HETEROGENEOUS DSS

A. Example & Proof of Theorem 1

We illustrate the proof of Th. 1 by an example for the special case in which the bandwidths depend only on the identity of the helper node. We compute the capacity of the DSS for this specific example, and show that it is strictly less than the upper bound of Th. 1. That is, it does not achieve the capacity of a homogenous system with the same average characteristics. More specifically, consider the heterogeneous DSS depicted in Fig. 1(a) with $(n, k, d) = (3, 2, 2)$ formed of 3 storage nodes v_1, v_2 and v_3 with storage capacities $(\alpha_1, \alpha_2, \alpha_3) = (1, 2, 2)$ and repair bandwidths $(\beta_1, \beta_2, \beta_3) = (1, 2, 2)$. The average node capacity is $\bar{\alpha} = 5/3$ and repair bandwidth is $\bar{\beta} = 10/3$. Th. 1 gives that the capacity of this DSS $C \leq 10/3 = 3.33$.

For this example, it is easy to see that the DSS capacity is $C = 3 \leq 10/3$. In fact, a user contacting nodes v_1 and v_2 cannot download more information than their total storage $\alpha_1 + \alpha_2 = 3$. This upper bound is achieved by the code in Fig. 1(a). The code stores a file of 3 units (x, y, z) in the

system. During repair the new node can download the whole file and store the lost piece of the data.

To obtain the upper bound in (4), we use the original heterogeneous DSS to construct a “bigger” homogeneous system. We obtain this new system by “glueing” together $n! = 3! = 6$ copies of the original DSS as shown in Fig. 1(b). Each copy corresponds to a different permutation of the nodes. In the figure, the i^{th} copy stores the file (x_i, y_i, z_i) . For example in Fig. 1(b), the first copy is the original system itself, the second corresponds to nodes v_1 and v_3 switching positions, and so on.

The “bigger” system is homogeneous because all its nodes have storage $\alpha = 10$ and repair bandwidth per helper $\beta = \gamma/d = 10$. The capacity C' of this system can be computed from (2): $C' = \sum_{i=1}^k \min \{ \alpha, (d-i+1) \frac{\gamma}{d} \} = 20$. Any scheme that can store a file of size C in the original DSS can be made into a scheme that can store a file of size $n!C$ in the “bigger” DSS. Therefore, we get $n!C \leq C'$ and $C \leq 10/3$.

We prove Th. 1 by generalizing the argument above. We start by describing formally the operation of adding, or combining, together multiple storage systems having the same number of nodes. Let DSS_1 and DSS_2 be two storage systems with nodes v_1^1, \dots, v_n^1 and v_1^2, \dots, v_n^2 , respectively. The new system that we refer to as DSS obtained by combining DSS_1 and DSS_2 is comprised of n nodes, say u_1, \dots, u_n . Node u_i has storage capacity $\alpha_i = \alpha_i^1 + \alpha_i^2$ (superscript $j, j = 1, 2$, denotes a parameter of system DSS_j). Moreover, when node u_j fails in DSS , the new node downloads $\beta_{ijs} = \beta_{ijs}^1 + \beta_{ijs}^2$ amount of information from helper node u_i (recall that S is the set of indices of the d helper nodes). We write $DSS = DSS_1 + DSS_2$.

Now, let DSS be the given heterogeneous system for which we wish to compute its capacity C . For each permutation $\sigma : \{1, \dots, n\} \rightarrow \{1, \dots, n\}$, we denote by DSS_σ the storage system with nodes $v_1^\sigma, \dots, v_n^\sigma$ such that $v_i^\sigma = v_{\sigma(i)}$. Let \mathcal{P}_n denote the set of all $n!$ permutations on the set $\{1, \dots, n\}$. We define a new “big” system by $DSS_b = \sum_{\sigma \in \mathcal{P}_n} DSS_\sigma$. The new system DSS_b is homogeneous with symmetric repair where the storage capacity per node α_b is given by $\alpha_b = (n-1)! \sum_{i=1}^n \alpha_i = n! \bar{\alpha}$, and the repair bandwidth per helper β_b is given by $\beta_b = n! \frac{\bar{\gamma}}{d}$. Therefore, the capacity C_b of DSS_b as given by (2) is $C_b = n! \sum_{i=1}^k \min \{ \bar{\alpha}, (d-i+1) \frac{\bar{\gamma}}{d} \}$. Any scheme achieving the capacity C of the original system can

be naturally extended to store a file of size $n!C$ in \mathcal{DSS}_b (see Fig. 1). Therefore, $C_b \geq n!C$, which gives the result of Th. 1.

Theorem 1 implies that symmetric repair, *i.e.*, downloading equal numbers of bits from each of the helpers, is optimal in a homogeneous system. To see this, consider a DSS with node storage capacity α , and a total repair bandwidth budget γ . A new node joining the system has the flexibility to arbitrarily split its repair bandwidth among the d helpers as long as the total amount of downloaded information does not exceed γ . Now, irrespective of how each new node splits its bandwidth budget, the average repair bandwidth in the system is the same, $\bar{\gamma} = \gamma$. If we apply Th. 1, we get an upper bound that matches exactly the capacity in (2) of a homogeneous DSS with symmetric repair. Hence, we obtain the result in Cor. 2.

B. Proof of Theorem 3

To avoid heavy notation, we focus on the case in which the repair bandwidth depends only on the helper node ($\beta_{ijS} = \beta_i$). We give in Th. 5 lower and upper bounds specific to this case. These bounds are similar to the ones in Th. 3, but can be tighter. The proof of Th. 3 follows the exact steps of the proof below and will be omitted here. Again, we index the nodes in increasing order of their capacities, $\alpha_1 \leq \alpha_2 \leq \dots \leq \alpha_n$. We also order the values of the repair bandwidths β to obtain the increasing sequence $\beta'_1 \leq \beta'_2 \leq \dots \leq \beta'_n$.

Theorem 5: The capacity C of a heterogeneous DSS, in which the repair bandwidth depends only on the identity of the helper node, is bounded as $C'_{\min} \leq C \leq C'_{\max}$, where

$$\begin{aligned} C'_{\min} &= \sum_{i=1}^k \min(\alpha_i, \beta'_1 + \beta'_2 + \dots + \beta'_{d-i+1}) \\ &= \min_{l=0, \dots, k} \left(\sum_{i=1}^l \alpha_i + \sum_{j=0}^{k-l-1} \sum_{i=1}^{d-l-j} \beta'_i \right), \text{ and} \\ C'_{\max} &= \sum_{i=1}^k \min(\alpha_i, \beta'_{i+1} + \beta'_{i+2} + \dots + \beta'_{d+1}) \\ &= \min_{l=0, \dots, k} \left(\sum_{i=1}^l \alpha_i + \sum_{j=1}^{k-l} \sum_{i=l+1+j}^{d+1} \beta'_i \right). \end{aligned} \quad (6)$$

The second expressions for C'_{\min} and C'_{\max} highlight the analogy with the bounds in Th. 3. Before proving Th. 5, we give some examples and discuss some special cases.

Example 6: Consider again the example in the previous section where $(n, k, d) = (3, 2, 2)$ and where the nodes parameters are $(\alpha_1, \beta_1) = (1, 1)$ and $(\alpha_2, \beta_2) = (\alpha_3, \beta_3) = (2, 2)$. Here, $C'_{\min} = 2$ and $C'_{\max} = 3$. Here, C'_{\max} is tighter than the average-based upper bound of Th. 1 which gives $C \leq 3.33$. Recall that the capacity for this system is $C = 3 = C'_{\max}$.

Example 7: Consider now a second DSS with $(n, k, d) = (3, 2, 2)$ and $(\alpha_1, \beta_1) = (5, 3), (\alpha_2, \beta_2) = (6, 4)$ and $(\alpha_3, \beta_3) = (7, 5)$. Here, $C'_{\min} = 9$ and $C'_{\max} = 11$, and Th. 1 gives $C \leq 10 < C'_{\max}$.

The upper and lower bounds can coincide in certain cases, which gives an exact expression for the capacity:

- 1) A homogeneous DSS, for which we recover the capacity expression in (2).

- 2) A DSS with uniform repair bandwidth, *i.e.*, $\beta_i = \beta, \forall i$. The capacity is $C = \sum_{i=1}^k \min(\alpha_i, (d-i+1)\beta)$.
- 3) Whenever $\alpha_i \leq \beta'_1, \forall i$. In this case, $C = \sum_{i=1}^k \alpha_i$.

To prove the upper and lower bounds in Th. 5, we first establish the following expression of the DSS capacity.

Theorem 8: The capacity C of a heterogeneous DSS is

$$C = \min_{\substack{(f_1, \dots, f_k) \\ f_i \neq f_j \text{ for } i \neq j}} \sum_{i=1}^k \min \left(\alpha_{f_i}, \min_{\substack{|S|=d+1-i \\ S \cap \{f_1, \dots, f_i\} = \emptyset}} \beta_S \right), \quad (8)$$

where for any $S \subseteq \{1, \dots, n\}$, $\beta_S = \sum_{i \in S} \beta_i$.

The proof of Th. 8 is a generalization of the proof in [1] of the capacity of a homogeneous system (2). Due to space limitations, we have omitted the proof here and it can be found in [15]. Next, we explain the intuition behind the proof. Consider the scenario depicted in Fig. 2 where nodes v_{f_1}, \dots, v_{f_k} fail and are repaired successively such that node v_{f_i} is repaired by downloading data from the previously repaired nodes $v_{f_1}, \dots, v_{f_{i-1}}$ and $d - (i - 1)$ other helper nodes in the system. Consider now a user contacting nodes v_{f_1}, \dots, v_{f_k} . The amount of “non-redundant” information that node v_{f_i} can give to the user is evidently limited by its storage capacity α_i on one hand, and on the other hand, by the amount of information β_{S_i} downloaded from the $d - i + 1$ helper nodes that are not connected to the user. Minimizing over all the choices of f_1, \dots, f_k gives the expression in (8).

It is not clear whether the capacity expression in (8) can be computed efficiently. For this reason we give upper and lower bounds that are easy to compute. To get the lower bound in (6), let $(f_1, \dots, f_k) = (f_1^*, \dots, f_k^*)$ be the minimizer of (8):

$$\begin{aligned} C &= \sum_{i=1}^k \min \left(\alpha_{f_i^*}, \min_{\substack{|S_i|=d+1-i \\ \{f_1^*, \dots, f_i^*\} \cap S_i = \emptyset}} \beta_{S_i} \right) \\ &\geq \sum_{i=1}^k \min(\alpha_{f_i^*}, \beta'_1 + \beta'_2 + \dots + \beta'_{d-i+1}) \\ &\geq \sum_{i=1}^{l^*} \alpha_i + \sum_{i=1}^{d-l^*} \beta'_i + \sum_{i=1}^{d-l^*-1} \beta'_i + \dots + \sum_{i=1}^{d-k+1} \beta'_i \\ &= \min_{l=0, \dots, k} \left(\sum_{i=1}^l \alpha_i + \sum_{j=0}^{k-l-1} \sum_{i=1}^{d-l-j} \beta'_i \right), \end{aligned} \quad (9)$$

where $l^*, 0 \leq l^* \leq k$ is the number of those cases in which $\alpha_{f_i^*}$ is smaller than or equal to the corresponding sum of β' 's.

The upper bound C'_{\max} is obtained by taking $(f_1, \dots, f_k) = (1, \dots, k)$ in (8) and following similar steps as above.

V. SECURITY

We now consider the case in which ℓ nodes in the system are compromised by a passive eavesdropper who can observe their downloaded and stored data, but cannot alter it. The *secrecy capacity* C_s of the system is the maximum amount of information that can be delivered to any user without revealing any information to the eavesdropper (perfect secrecy).




Fig. 2. A series of k failures and repairs in the DSS that explains the capacity expression in (8). Nodes v_{f_1}, \dots, v_{f_k} fail successively and are repaired as depicted above. The amount of “new” information that node v_{f_i} can give the user is the minimum between his capacity α_{f_i} and downloaded data β_{S_i} .

Formally, let S be the information source that represents the file stored on the DSS. A user contacts the nodes in any set $B \subset \{v_1, \dots, v_n\}$ of size k and downloads their stored data denoted by C_B . The user should be able to decode the file, which implies $H(S|C_B) = 0$. Let E be the set of the ℓ compromised nodes, and D_E be the data observed by the eavesdropper. The perfect secrecy condition implies that $H(S|D_E) = H(S)$. Following the definition in [14], we write the secrecy capacity as

$$C_s(\alpha, \gamma) = \sup_{\substack{H(S|C_B)=0 \forall B \\ H(S|D_E)=H(S) \forall E}} H(S). \quad (10)$$

Finding the secrecy capacity of a DSS is a hard problem and is still open in general, even for the class of homogeneous systems. Let $C_s^{ho}(\alpha, \beta, \ell)$ denote the secrecy capacity of a homogeneous DSS implementing symmetric repair and having ℓ compromised nodes. Following the same steps as in the proof of Th. 1, we can show that the secrecy capacity C_s of a heterogeneous DSS cannot exceed that of a homogeneous DSS having the same average resources.

Theorem 9: Consider a heterogeneous DSS with average storage capacity per node $\bar{\alpha}$, average repair bandwidth $\bar{\gamma}$, and ℓ compromised nodes. The secrecy capacity of this system is upper bounded by

$$C_s \leq C_s^{ho}(\bar{\alpha}, \bar{\gamma}, \ell). \quad (11)$$

Eqs. (11) and (3) imply the following upper bound in Th. 4:

$$C_s \leq \sum_{i=\ell+1}^k \min \left\{ \bar{\alpha}, (d-i+1) \frac{\bar{\gamma}}{d} \right\}. \quad (12)$$

Using Th. 9, we deduce that symmetric repair is also optimal here and maximizes the secrecy capacity of a DSS.

Corollary 10: Symmetric repair maximizes the secrecy capacity of a homogeneous system with a given budget on total repair bandwidth.

VI. CONCLUSION

We have studied heterogeneous distributed storage systems. Nodes in these systems can have different storage capacities and repair bandwidths. We have focused on determining the information theoretic capacity of these systems, *i.e.*, the maximum amount of information they can store under a reliability

constraint. We have proved an upper bound on the capacity that depends on the average node resources. Moreover, we have given an expression for the system capacity when the nodes’ parameters are known. This expression may be hard to compute, but we use it to derive additional upper and lower bounds that are easy to evaluate. We have also studied the case in which the system is compromised by an eavesdropper, and have provided bounds on the system secrecy capacity under a perfect secrecy constraint. Our results imply that symmetric repair maximizes the capacity of a homogeneous system, which justifies the repair model used in the literature. Problems that remain open include finding an efficient algorithm to compute the capacity of a heterogeneous distributed storage system, as well as efficient code constructions.

ACKNOWLEDGEMENTS

This research was supported in part by the U. S. Army Research Office under MURI Grant W911NF-11-1-0036, in part by the U. S. National Science Foundation under Grant CCF-1016671, and in part by the Academy of Finland under Grant #131745. This research was partly carried out under the COST Action IC1104.

REFERENCES

- [1] A. Dimakis, P. Godfrey, Y. Wu, M. Wainright, and K. Ramchandran, “Network coding for distributed storage systems,” *IEEE Transactions on Information Theory*, vol. 56, pp. 4539–4551, 2010.
- [2] S. Ghemawat, H. Gobioff, and S.-T. Leung, “The Google file system,” in *Proc. 19th ACM Symposium on Operating Systems Principles*, 2003.
- [3] A. G. Dimakis, K. Ramchandran, Y. Wu, and C. Suh, “A survey on network codes for distributed storage,” *arXiv:1004.4438*, 2010.
- [4] J. Kubiatowicz, D. Bindel, Y. Chen, S. Czerwinski, P. Eaton, D. Geels, R. Gummadi, S. Shea, H. Weatherspoon, W. Weimer, C. Wells, and B. Zhao, “Oceanstore: An architecture for global-scale persistent storage,” in *Proc. 9th International Conference on Architectural Support for Programming Languages and Operating Systems*, 2000.
- [5] “<http://techcrunch.com/2012/07/11/space-monkey-seed-round/>.”
- [6] H. Zhang, M. Chen, A. Parek, and K. Ramchandran, “A distributed multichannel demand-adaptive p2p vod system with optimized caching and neighbor-selection,” in *Proc. SPIE*, 2011.
- [7] S. Pawar, S. El Rouayheb, H. Zhang, K. Lee, and K. Ramchandran, “Codes for a distributed caching based video-on-demand system,” in *Proc. Asilomar Conference on Signals, Systems and Computers*, 2011.
- [8] N. Golrezaei, A. G. Dimakis, and A. F. Molisch, “Wireless device-to-device communications with distributed caching,” in *Proc. IEEE International Symposium on Information Theory*, 2012.
- [9] Y. Wu, “A construction of systematic MDS codes with minimum repair bandwidth,” *arXiv:0910.2486*, 2009.
- [10] V. T. Van, C. Yuen, and J. Li, “Non-homogeneous distributed storage systems,” in *arXiv:1208.2078*, 2012.
- [11] D. Leong, A. G. Dimakis, and T. Ho, “Distributed storage allocations,” in *arXiv:1011.5287v3*, 2012.
- [12] V. Ntranos, G. Caire, and A. G. Dimakis, “Allocations for heterogeneous distributed storage,” in *Proc. IEEE International Symposium on Information Theory*, 2012.
- [13] N. B. Shah, K. V. Rashmi, and V. Kumar, “A flexible class of regenerating codes for distributed storage,” in *Proc. IEEE International Symposium on Information Theory*, 2010.
- [14] S. Pawar, S. El Rouayheb, and K. Ramchandran, “Securing dynamic distributed storage systems against eavesdropping and adversarial attacks,” *IEEE Transactions on Information Theory*, vol. 58, pp. 6734–6753, 2012.
- [15] T. Ernvall, S. El Rouayheb, C. Hollanti, H. V. Poor, “Capacity and Security of Heterogeneous Distributed Storage Systems (extended version),” *arXiv:1211.0415*, 2012.